

# **TTS 11.0 COOKBOOK**

## **(NSD NOSQL DAY02)**

版本编号 11.0

2019-06

达内 IT 培训集团

## NSD NOSQL DAY02

### 1. 案例 1：部署 redis 集群

- 问题

- 具体要求如下：
- 准备集群环境
- 安装 redis 并创建集群
- 查看集群信息

- 方案

搭建 redis 集群，拓扑规划如图-1 所示：

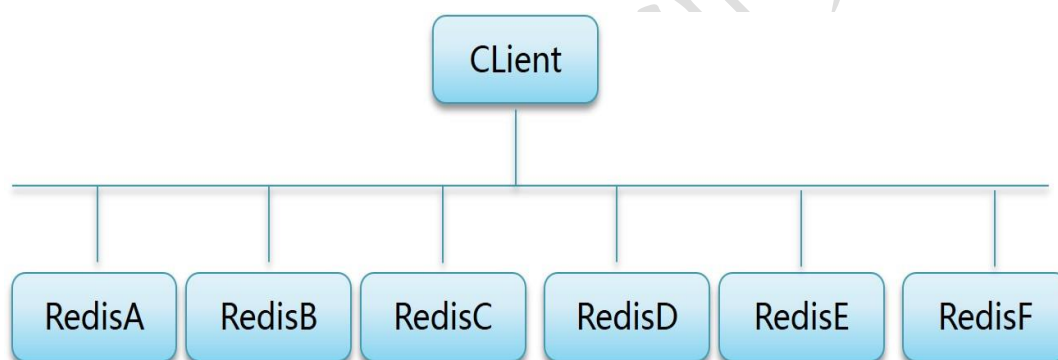


图 - 1

IP，端口规划如表-1 所示：

表-1

主机名	IP 地址	端口号
redisA	192.168.4.51	6351
redisB	192.168.4.52	6352
redisC	192.168.4.53	6353
redisD	192.168.4.54	6354
redisE	192.168.4.55	6355
redisF	192.168.4.56	6356

- 步骤

实现此案例需要按照如下步骤进行。

### 步骤一：准备集群环境

- 1) 按照表-1 配置主机名, ip 地址, 配置 yum 源 (系统源) 这里不再操作
- 2) 把 redis 的软件包传到 6 台数据库服务器上面, 安装 redis 服务器, 六台服务器同样操作 (以 51 为例)

```
[root@redisA ~]# yum -y install gcc gcc-c++ make
[root@redisA ~]# cd redis
redis/      redis-cluster/
[root@redisA ~]# cd redis/
[root@redisA redis]# ls
lnmp  redis-4.0.8.tar.gz
[root@redisA redis]# tar -xf redis-4.0.8.tar.gz
[root@redisA redis]# cd redis-4.0.8/
[root@redisA redis-4.0.8]# make && make install
[root@redisA redis-4.0.8]# ./utils/install_server.sh
Welcome to the redis service installer
This script will help you easily set up a running redis server

Please select the redis port for this instance: [6379]
Selecting default: 6379
Please select the redis config file name [/etc/redis/6379.conf]
Selected default - /etc/redis/6379.conf
Please select the redis log file name [/var/log/redis_6379.log]
Selected default - /var/log/redis_6379.log
Please select the data directory for this instance [/var/lib/redis/6379]
Selected default - /var/lib/redis/6379
Please select the redis executable path [/usr/local/bin/redis-server]
Selected config:
Port          : 6379
Config file   : /etc/redis/6379.conf
Log file      : /var/log/redis_6379.log
Data dir      : /var/lib/redis/6379
Executable    : /usr/local/bin/redis-server
Cli Executable : /usr/local/bin/redis-cli
Is this ok? Then press ENTER to go on or Ctrl-C to abort.
Copied /tmp/6379.conf => /etc/init.d/redis_6379
Installing service...

Successfully added to chkconfig!
Successfully added to runlevels 345!
Starting Redis server...
Installation successful!           //安装成功

[root@redisA redis-4.0.8]# ss -antlp | grep 6379           //查看时有端口
LISTEN      0          128        127.0.0.1:6379
users:(("redis-server",pid=10788,fd=6))                  *:*
```

- 2) 修改配置文件, 6 台 redis 服务器都要修改 (以 51 为例子)

```
[root@redisA redis-4.0.8]# /etc/init.d/redis_6379 stop
//停止已经开启的 redis 服务
Stopping ...
Waiting for Redis to shutdown ...
Redis stopped
```

```
[root@redisA redis-4.0.8]# vim /etc/redis/6379.conf
...
bind 192.168.4.51      //修改 ip
port 6351              //不允许相同，只指定物理接口的 ip
daemonize yes          //以守护进程方式运行
pidfile /var/run/redis_6351.pid
cluster-enabled yes    //是否启用集群，前提是以守护进程方式运行
cluster-config-file nodes-6351.conf
//存储集群信息的配置文件，自动生成，不允许相同
cluster-node-timeout 5000 //集群节点通信超时时间
...
[root@redisA redis-4.0.8]# /etc/init.d/redis_6379 start //启动服务
Starting Redis server...
[root@redisA redis-4.0.8]# ss -antlp | grep 6351 //查看有端口
LISTEN      0            128          192.168.4.51:6351          *:*
users:((("redis-server",pid=11092,fd=6))
LISTEN      0            128          192.168.4.51:16351        *:*
users:((("redis-server",pid=11092,fd=8)) //16051: 集群中的主机通信时用的端口
[root@redisA redis-4.0.8]# ps -C redis
PID TTY          TIME CMD
```

**注意：其他几台主机在修改时请注意 ip，端口等的修改，不要和 51 主机的一样**

### 3) 关闭防火墙 51-56 主机 (以 51 为例子)

```
[root@redisA redis-4.0.8]# getenforce
Permissive

[root@redisA redis-4.0.8]# systemctl disable firewalld
//关闭防火墙不自启
```

### 4) 查看集群信息

```
[root@redisA redis-4.0.8]# redis-cli -h 192.168.4.51 -p 6351
192.168.4.51:6351> ping
PONG
192.168.4.51:6351> cluster info
cluster_state:fail
cluster_slots_assigned:0
cluster_slots_ok:0
cluster_slots_pfail:0
cluster_slots_fail:0
cluster_known_nodes:1
cluster_size:0
...

192.168.4.51:6351> cluster nodes
f81f997d5ed988ec1587558e78d5f7dbc96abcbf :6351@16351 myself,master - 0 0 0 connected
```

## 步骤二：创建集群 (在任意一台上执行创建集群的脚本都可以) 这里在 51 上面执行

### 1) 部署 ruby 脚本运行环境 (在 51 上面执行)

```
[root@redisA redis-4.0.8]# cd /root/redis-cluster/
[root@redisA redis-cluster]# ls
redis-3.2.1.gem  ruby-devel-2.0.0.648-30.el7.x86_64.rpm
```

```
[root@redisA redis-cluster]# yum -y install ruby rubygems

[root@redisA redis-cluster]# rpm -ivh --nodeps \
  ruby-devel-2.0.0.648-30.el7.x86_64.rpm
warning: ruby-devel-2.0.0.648-30.el7.x86_64.rpm: Header V3 RSA/SHA256 Signature,
key ID f4a80eb5: NOKEY
Preparing... ##### [100%]
Updating / installing...
  1:ruby-devel-2.0.0.648-30.el7 ##### [100%]
[root@redisA redis-cluster]# which gem
/usr/bin/gem
[root@redisA redis-cluster]# gem install redis
Successfully installed redis-3.2.1
Parsing documentation for redis-3.2.1
Installing ri documentation for redis-3.2.1
1 gem installed
```

## 2) 生成创建集群的脚本

```
[root@redisA redis-cluster]# cd /root/redis/redis-4.0.8/src/
[root@redisA src]# cp redis-trib.rb /usr/local/bin/
[root@redisA src]# ll /usr/local/bin/redis-trib.rb
-rwxr-xr-x. 1 root root 65991 Sep 27 16:12 /usr/local/bin/redis-trib.rb
```

## 3) 创建集群

```
[root@redisA src]# redis-trib.rb create --replicas 1 \
192.168.4.51:6351 192.168.4.52:6352 \
192.168.4.53:6353 192.168.4.54:6354 \
192.168.4.55:6355 192.168.4.56:6356
//--replicas 1 给每一个主配置一个从库
[root@redisA log]# redis-trib.rb create --replicas 1 192.168.4.51:6351
192.168.4.52:6352 192.168.4.53:6353 192.168.4.54:6354 192.168.4.55:6355
192.168.4.56:6356
>>> Creating cluster
>>> Performing hash slots allocation on 6 nodes...
Using 3 masters:
192.168.4.51:6351
192.168.4.52:6352
192.168.4.53:6353
Adding replica 192.168.4.55:6355 to 192.168.4.51:6351
Adding replica 192.168.4.56:6356 to 192.168.4.52:6352
Adding replica 192.168.4.54:6354 to 192.168.4.53:6353
...
...
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
```

## 4) 查看集群信息，任意一台主机访问本机的 redis 服务查看即可

```
cluster info    查看集群信息
cluster nodes   查看集群节点信息
```

```
[root@redisA log]# redis-cli -h 192.168.4.52 -p 6352
192.168.4.52:6352> CLUSTER INFO
cluster_state:ok           //状态
cluster_slots_assigned:16384
cluster_slots_ok:16384
```

```
cluster_slots_pfail:0
cluster_slots_fail:0
cluster_known_nodes:6
cluster_size:3
cluster_current_epoch:6
cluster_my_epoch:2
cluster_stats_messages_ping_sent:367
cluster_stats_messages_pong_sent:327
cluster_stats_messages_meet_sent:5
cluster_stats_messages_sent:699
cluster_stats_messages_ping_received:327
cluster_stats_messages_pong_received:372
cluster_stats_messages_received:699
```

```
192.168.4.52:6352> CLUSTER NODES //查看集群节点信息
63afbb5e7d63b1f142358634578a3488e3c9e634 192.168.4.54:6354@16354 slave
bc5c4e082a5a3391b634cf433a6486c867cfc44b 0 1538039278871 4 connected
bc5c4e082a5a3391b634cf433a6486c867cfc44b 192.168.4.53:6353@16353 master - 0
1538039278571 3 connected 10923-16383
28e06c5f24a2b6c6412f81369e09bc9653cc51ff 192.168.4.56:6356@16356 slave
8568fbd73cb296cad6915d524e34761b2114af47 0 1538039278069 6 connected
7e8d9121f44d8331ff55b45c218b87df9bda1b70 192.168.4.55:6355@16355 slave
a3083123bc5c87a76aab2655171634d4ee84f418 0 1538039278000 5 connected
8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352@16352 myself,master - 0
1538039277000 2 connected 5461-10922
a3083123bc5c87a76aab2655171634d4ee84f418 192.168.4.51:6351@16351 master - 0
1538039277869 1 connected 0-5460
192.168.4.52:6352>
```

## 5) 测试集群

命令:

```
redis-cli -c -h ip地址 -p 端口
```

```
[root@redisA log]# redis-cli -c -h 192.168.4.51 -p 6351
192.168.4.51:6351> set name jim
-> Redirected to slot [5798] located at 192.168.4.52:6352
OK
192.168.4.52:6352> get name
"jim"
192.168.4.52:6352> set class linux
OK
192.168.4.52:6352> get class
"linux"
192.168.4.52:6352> set pay 26800
-> Redirected to slot [4013] located at 192.168.4.51:6351
OK
192.168.4.51:6351> get pay
"26800"
```

集群不能用的情况:

有半数或者半数以上的主库机器挂掉, 集群就不能用了

把一个从库升级成主, 没有从库, 集群不能用 (前提是: 有半数或者半数以上的主库机器挂掉)

一个主库挂掉, 它的从库自动顶替为主库, 正常使用 (前提是: 有半数或者半数以上的主库机器能用), 挂掉的主库修复好后, 会成为从库, 不会抢占为主

## 6) 集群节点选举策略 (三主, 三从)

停止某个主库的 redis 服务, 对应的从库会自动升级为主库

先查看节点信息的主从情况

```
[root@redisA log]# redis-cli -c -h 192.168.4.51 -p 6351
192.168.4.51:6351> CLUSTER nodes
...
8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352@16352 master - 0
1538040400840 2 connected 5461-10922
28e06c5f24a2b6c6412f81369e09bc9653cc51ff 192.168.4.56:6356@16356 slave
8568fbd73cb296cad6915d524e34761b2114af47 0 1538040400000 6 connected
...
192.168.4.51:6351>
```

看谁是谁的从库, 如:

看节点前后的编号 id 是否有相同的

如: 8568fbd73cb296cad6915d524e34761b2114af47

发现 52 的从库为 56

停止主库 52

```
[root@redisA log]# redis-cli -h 192.168.4.52 -p 6352 shutdown
[root@redisA log]# redis-cli -c -h 192.168.4.51 -p 6351
192.168.4.51:6351> CLUSTER NODES
...
8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352@16352 master,fail -
1538041052349 1538041051000 2 disconnected //52 的主库坏掉
28e06c5f24a2b6c6412f81369e09bc9653cc51ff 192.168.4.56:6356@16356 master - 0
1538041066000 7 connected 5461-10922 //56 成为主库
...

```

开启 52, 发现 52 成为从库

```
[root@redisB redis-4.0.8]# /etc/init.d/redis_6352 start
Starting Redis server...
[root@redisA log]# redis-cli -c -h 192.168.4.51 -p 6351
192.168.4.51:6351> CLUSTER NODES
8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352@16352 slave
28e06c5f24a2b6c6412f81369e09bc9653cc51ff 0 1538041254000 7 connected

```

## 2. 案例 2: 管理 redis 集群

### • 问题

- 具体要求如下:
- 练习添加主机
- 练习删除主机

### • 步骤

实现此案例需要按照如下步骤进行。

### 步骤一：添加主机

1) 部署一台新 redis 服务器，ip 为 192.168.4.58，装包，初始化，启用集群配置，重启服务（这里之前已经操作过，不会的可以参考案例 1）

2) 添加集群 4.58（添加 master 节点）

格式：redis-trib.rb 选项 参数

选项：add-node 添加主机（不指定角色为主）

由于之前是在 51 上面创建 ruby 脚本，所以只有 51 上面有 redis-trib.rb 命令，在 51 上面执行

```
[root@redisA ~]# redis-trib.rb add-node 192.168.4.58:6358 192.168.4.51:6351
>>> Adding node 192.168.4.58:6358 to cluster 192.168.4.51:6351
>>> Performing Cluster Check (using node 192.168.4.51:6351)
S: a3083123bc5c87a76aab2655171634d4ee84f418 192.168.4.51:6351
   slots: (0 slots) slave
   replicates 7e8d9121f44d8331ff55b45c218b87df9bda1b70
M: 7e8d9121f44d8331ff55b45c218b87df9bda1b70 192.168.4.55:6355
   slots:0-5460 (5461 slots) master
   1 additional replica(s)
S: 8568fbd73cb296cad6915d524e34761b2114af47 192.168.4.52:6352
   slots: (0 slots) slave
   replicates 28e06c5f24a2b6c6412f81369e09bc9653cc51ff
M: bc5c4e082a5a3391b634cf433a6486c867cfc44b 192.168.4.53:6353
   slots:10923-16383 (5461 slots) master
   1 additional replica(s)
S: 63afbb5e7d63b1f142358634578a3488e3c9e634 192.168.4.54:6354
   slots: (0 slots) slave
   replicates bc5c4e082a5a3391b634cf433a6486c867cfc44b
M: 28e06c5f24a2b6c6412f81369e09bc9653cc51ff 192.168.4.56:6356
   slots:5461-10922 (5462 slots) master
   1 additional replica(s)
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
>>> Check slots coverage...
[OK] All 16384 slots covered.
>>> Send CLUSTER MEET to node 192.168.4.58:6358 to make it join the cluster.
[OK] New node added correctly.
```

3) 检查集群主机的状态信息

选项：check 检查权限

```
[root@redisA ~]# redis-trib.rb check 192.168.4.58:6358 //查看状态
>>> Performing Cluster Check (using node 192.168.4.58:6358)
M: c5e0da48f335c46a2ec199faa99b830f537dd8a0 192.168.4.58:6358
   slots: (0 slots) master //发现没有 hash 槽
   0 additional replica(s)
M: 7e8d9121f44d8331ff55b45c218b87df9bda1b70 192.168.4.55:6355
   slots:0-5460 (5461 slots) master
   1 additional replica(s)
...
S: a3083123bc5c87a76aab2655171634d4ee84f418 192.168.4.51:6351
   slots: (0 slots) slave
   replicates 7e8d9121f44d8331ff55b45c218b87df9bda1b70
[OK] All nodes agree about slots configuration.
>>> Check for open slots...
```



```
>>> Check slots coverage...  
[OK] All 16384 slots covered.
```

#### 4) 手动对集群进行分片迁移

选项: reshard 重新分配 hash 槽

```
[root@redisA ~]# redis-trib.rb reshard 192.168.4.58:6358  
How many slots do you want to move (from 1 to 16384)?4096  
//拿出多少个 hash 槽给主机 192.168.4.58  
What is the receiving node ID? c5e0da48f335c46a2ec199faa99b830f537dd8a0  
//主机 192.168.4.58 的 id 值  
Source node #1:all //从当前所有的主里面获取 hash 槽  
Do you want to proceed with the proposed reshard plan (yes/no)?yes  
...  
Moving slot 12283 from 192.168.4.53:6353 to 192.168.4.58:6358:  
Moving slot 12284 from 192.168.4.53:6353 to 192.168.4.58:6358:  
Moving slot 12285 from 192.168.4.53:6353 to 192.168.4.58:6358:  
Moving slot 12286 from 192.168.4.53:6353 to 192.168.4.58:6358:  
Moving slot 12287 from 192.168.4.53:6353 to 192.168.4.58:6358:
```

再次查看发现 4.58 有 4096 个 hash slot

```
[root@redisA ~]# redis-trib.rb check 192.168.4.58:6358  
>>> Performing Cluster Check (using node 192.168.4.58:6358)  
M: c5e0da48f335c46a2ec199faa99b830f537dd8a0 192.168.4.58:6358  
slots:0-1364,5461-6826,10923-12287 (4096 slots) master  
0 additional replica(s)
```

#### 5) 删除 master 角色的主机

先删除主机占用的 hash 槽

```
[root@redisA ~]# redis-trib.rb reshard 192.168.4.58:6358  
How many slots do you want to move (from 1 to 16384)?4096  
//移除 hash 槽的个数  
What is the receiving node ID? bc5c4e082a5a3391b634cf433a6486c867cfc44b  
//要移动给谁的 id 即目标主机 (这里可以随机写一个 master 的 ID)  
Source node #1: c5e0da48f335c46a2ec199faa99b830f537dd8a0  
//从谁那移动即源主机 (这里写 4.58 的 ID)  
Source node #2:done //设置完毕  
...  
Moving slot 12282 from c5e0da48f335c46a2ec199faa99b830f537dd8a0  
Moving slot 12283 from c5e0da48f335c46a2ec199faa99b830f537dd8a0  
Moving slot 12284 from c5e0da48f335c46a2ec199faa99b830f537dd8a0  
Moving slot 12285 from c5e0da48f335c46a2ec199faa99b830f537dd8a0  
Moving slot 12286 from c5e0da48f335c46a2ec199faa99b830f537dd8a0  
Moving slot 12287 from c5e0da48f335c46a2ec199faa99b830f537dd8a0  
Do you want to proceed with the proposed reshard plan (yes/no)?yes //提交  
...  
Moving slot 12282 from 192.168.4.58:6358 to 192.168.4.53:6353:  
Moving slot 12283 from 192.168.4.58:6358 to 192.168.4.53:6353:  
Moving slot 12284 from 192.168.4.58:6358 to 192.168.4.53:6353:  
Moving slot 12285 from 192.168.4.58:6358 to 192.168.4.53:6353:  
Moving slot 12286 from 192.168.4.58:6358 to 192.168.4.53:6353:  
Moving slot 12287 from 192.168.4.58:6358 to 192.168.4.53:6353:
```

删除集群主机 4.58(删除之后 redis 服务自动关闭)

```
[root@redisA ~]# redis-trib.rb del-node 192.168.4.58:6358 \  
c5e0da48f335c46a2ec199faa99b830f537dd8a0 //删除谁+删除的 id  
>>> Removing node c5e0da48f335c46a2ec199faa99b830f537dd8a0 from cluster  
192.168.4.58:6358  
>>> Sending CLUSTER FORGET messages to the cluster...  
>>> SHUTDOWN the node.
```

#### 6) 添加从节点主机, 随机添加

```
[root@redisA ~]# redis-trib.rb add-node --slave \  
192.168.4.57:6357 192.168.4.51:6351  
>>> Adding node 192.168.4.57:6357 to cluster 192.168.4.51:6351  
>>> Performing Cluster Check (using node 192.168.4.51:6351)  
.....  
.....  
[OK] All 16384 slots covered.  
Automatically selected master 192.168.4.51:6351  
>>> Send CLUSTER MEET to node 192.168.4.57:6357 to make it join the cluster.  
Waiting for the cluster to join.  
>>> Configure node as replica of 192.168.4.51:6351.  
[OK] New node added correctly.
```

#### 7) 移除从节点, 从节点主机没有槽位范围, 直接移除即可

命令格式:

redis-trib.rb del-node 192.168.4.57:6357 主机 id 值

```
[root@redisA ~]# redis-trib.rb del-node 192.168.4.57:6357 \  
f6649ea99b2f01faca26217691222c17a3854381  
>>> Removing node f6649ea99b2f01faca26217691222c17a3854381  
from cluster 192.168.4.57:6351  
>>> Sending CLUSTER FORGET messages to the cluster...  
>>> SHUTDOWN the node.
```